

## INTRODUCTION

.....  .....

### The Lay of the Land

There is so much to know about our world. And for those who are the least bit curious, we have more resources than ever to give us the insights we seek. We can turn to a variety of scientists, doctors, economists, historians, and journalists to help us better understand ourselves, our world, and our place within it.

But there is a set of vital questions that such experts will never answer. These are questions about how we ought to live. Sure, financial advisors can tell us how we ought to invest our money. Personal trainers can advise us on getting in shape. Career counselors can steer us in one direction or another. But if we are interested instead in what our guiding ideals should be, in what sort of life is worth living, in how we should treat one another, then we must turn to philosophy. *Ethics*—also known as *moral philosophy*—is the branch of knowledge concerned with answering such questions.

The field of ethics is vast, and—bad news first—there is no chance of covering all of its interesting and important issues within these pages. In selecting the topics for treatment, I have chosen those that seem to me most central. These can be grouped under three headings, each representing a core area of moral philosophy:

1. **Value theory**<sup>1</sup>: What is the good life? What is worth pursuing for its own sake? How do we improve our lot in life?
2. **Normative ethics**: What are our fundamental moral duties? Which character traits count as virtues, which as vices, and why? Who should our role models be? Do the ends always justify the means, or are there certain types of action that should never be done under any circumstances?
3. **Metaethics**: What is the status of moral claims and advice? Can ethical theories, moral principles, or specific moral verdicts be true? If so, what makes them true? Can we gain moral wisdom? If so, how? Do we always have good reason to do our moral duty?

The structure of this book mirrors this threefold division. The first part is focused on the good life, with an emphasis on explaining the nature and sources of well-being. We ask, for instance, about whether happiness is the be-all and end-all of a good life, the only thing desirable for its own sake. And, naturally, we'll consider views that deny this, including, most importantly, the theory that tells us that getting what we want—whatever we want—is the key to the good life.

Then it's off to normative ethics, which is devoted to examining our moral relations with one another. Who counts—are animals, ecosystems, or fetuses morally important in their own right? Is there a fundamental moral rule, such as the **golden rule**, that can justify all of our specific moral duties? What role do virtue, self-interest, and justice play in morality? Are we ever allowed to break the moral rules? If so, when and why? These are among the most important questions taken up in normative ethics.

Finally, to metaethics. This part of moral philosophy asks questions about the other two. Specifically, it asks about the status of ethical claims, rather than about their content. We all have views about what is right and good. Are these just matters of taste? Is moral authority based on personal approval? Social customs? God's commands? Or none of the above? Is morality in more or less good working order, or is it just a convenient fiction that keeps us in our place? These are the questions that we will take up in the last section of the book.

---

1. All technical terms and phrases that appear in **boldface** are defined in the Glossary at the end of the book.

There is no shortage of folks offering advice about these questions. The self-help industry has its gurus, motivational speakers, and best sellers, each aimed at guiding us on the path to a good life. Political pundits, religious leaders, and editorial writers are more than happy to offer us their blueprints for righteous living. They don't always agree, of course. It would be nice to have a way to sort out the decent advice from the rest.

Those of you turning to philosophical ethics for the first time are likely to be hoping for something that I can't provide, namely, a simple recipe for doing the sorting. It is perfectly natural to want a clear method for distinguishing correct from incorrect answers about the good life and our moral duty. Indeed, when I first went to college, I enrolled in a philosophy course hoping for just such a thing. My failure to find it left me deeply disappointed. I abandoned philosophy for a few years, and even dropped out of college for a while. After I returned, I went looking for it again. I've finally realized that in this area of life, while there is plenty of good advice, it can't be summed up in one snappy formula, captured in a neat slogan that can be inscribed in a fortune cookie or on a bumper sticker.

Ethics is hard. It needn't be weakness or fuzzy thinking that stands in the way of knowing the right thing to do, or the proper goals to strive for. We are right to be puzzled by the moral complexity we find in our lives. While we might yearn for clarity and simplicity, this wish for easy answers is bound to be repeatedly frustrated.

## Skepticism about Ethics

When people learn of the difficulties that face each important attempt to solve ethical puzzles, they often give in to skepticism. The major temptation is to regard the entire enterprise as bankrupt, or to think that all ethical views are equally plausible.

Doubts about morality are plentiful, and it would be silly to ignore them in a book that is so focused on trying to improve our moral understanding. Chapters 19, 20, and 21 are entirely devoted to such doubts; those who feel them acutely might do best to start with those chapters, and then work their way to the other parts of the book that are focused on the good life and normative ethics.

For now, let me say just a few things to the doubters. Perhaps the most important is this: among those who have thought longest and hardest about ethics, the view that morality is all make-believe, or that all moral

standards are correct only relative to individuals or societies, is deeply controversial. There are *lots* of problems with such views. Some of these problems may be devastating.

As a result, it would be a serious mistake just to assume that morality is a fiction, or that personal or cultural opinion is the ultimate measure of what is right and wrong. We must follow the arguments where they lead. They may indeed lead us ultimately to embrace such positions. But they might not. And we can't know one way or the other until we've actually done the hard work.

Since I really love this part of ethics—the metaethical part—I can't resist saying just a bit more here. In my experience, most of those who harbor serious doubts about morality base their skepticism on one or more of the following considerations:

- (A) Individuals constantly disagree about what's right and wrong, and societies do, too. If there were some objective truth in ethics, then we should expect all really smart people to agree on it. They don't. So there is no objective truth in ethics.
- (B) There are universally correct moral standards only if God exists. But God doesn't exist, so ethics is just a "human construct."
- (C) Science tells us the truth about the world, and science says nothing about what's right and wrong. And that's because nothing really is right or wrong.
- (D) If there were a universal ethic, then that would make it okay for some people to impose their own views on others. But that's not okay at all. Therefore there is no universal ethic.
- (E) If there were objective moral rules, then it would always be wrong to break them. But every rule admits of exceptions; no moral rule is absolute. That shows that we do make up the moral rules after all.

This is going to sound like cheating, but here goes: every single one of these arguments is problematic. I'm not going to defend that claim right now—that's what the last three chapters of this book are for. So I don't expect you to believe me (yet). Still, there is a lesson here: until these (or other) arguments are laid out with care and successfully defended, we are in no position to assume that the skeptics about morality are right.

I think you'll soon see that we can make a lot of progress in our moral thinking. And even if morality *is* in some way a human invention, there is still lots to learn, and there are many ways to make mistakes when

thinking about what is good and right. It's important to avoid these errors. Doing moral philosophy can help with this.

Look at it this way. Lots of people believe that when it comes to art, beauty is in the eye of the beholder—there are no objective, universal standards of good taste. Suppose that's true. And suppose that morality is just like art in this respect. Still, our tastes can be educated and improved. Many people are much wiser than I am about music and painting, for instance. Even if there are no universal standards of good taste, it would be silly of me to pass up a chance to talk with people who have thought long and hard about artistic matters. Why should I dismiss their opinions and refuse to hear them out? I'm no genius. Maybe I could learn a thing or two.

That's *exactly* the right attitude to take about ethics. Especially when so much is at stake—the very quality of our life and our relations with others—it would be terrible to close our minds to new and challenging ideas. Those who have thought so hard about the central questions of existence may well have something to teach us.

I encourage you to resist the diagnosis that in ethics, anything goes. As you'll see, good moral thinking is disciplined thinking. There are many ways that we can go wrong in our moral reflections, and failure here can have the most disastrous results. Though it is sometimes hard to know when we have got it right in ethics, it is often very easy to know when we (or others) have made a mistake. There are clear cases of people ruining their lives, or doing morally horrific things. We should keep that in mind before siding too quickly with a skepticism that says that every moral view is as good as every other.

## Ethical Starting Points

One of the puzzles about moral thinking is knowing where to begin. Some skeptics about morality deny that there are any proper starting points for ethical reflection. They believe that moral reasoning is simply a way of rationalizing our biases and gut feelings. This outlook encourages us to be lax in moral argument and, worse, supports an attitude that no moral views are any better than others. While this sort of skepticism might be true, we shouldn't regard it as the default view of ethics. We should accept it only as a last resort.

In the meantime, let's consider some fairly plausible ethical assumptions, claims that can get us started in our moral thinking. The point of the

exercise is to soften you up to the idea that we are not just spinning our wheels when thinking morally. There are reasonable constraints that can guide us when thinking about how to live. Here are some of them:

- *Neither the law nor tradition is immune from moral criticism.* The law does not have the final word on what is right and wrong. Neither does tradition. Actions that are legal, or customary, are sometimes morally mistaken.
- *Everyone is morally fallible.* Everyone has some mistaken ethical views, and no human being is wholly wise when it comes to moral matters.
- *Friendship is valuable.* Having friends is a good thing. Friendships add value to your life. You are better off when there are people you care deeply about, and who care deeply about you.
- *We are not obligated to do the impossible.* Morality can demand only so much of us. Moral standards that are impossible to meet are illegitimate. Morality must respect our limitations.
- *Children bear less moral responsibility than adults.* Moral responsibility assumes an ability on our part to understand options, to make decisions in an informed way, and to let our decisions guide our behavior. The fewer of these abilities you have, the less blameworthy you are for any harm you might cause.
- *Justice is a very important moral good.* Any moral theory that treats justice as irrelevant is deeply suspect. It is important that we get what we deserve, and that we are treated fairly.
- *Deliberately hurting other people requires justification.* The default position in ethics is this: do no harm. It is sometimes morally acceptable to harm others, but there must be an excellent reason for doing so or else the harmful behavior is unjustified.
- *Equals ought to be treated equally.* People who are alike in all relevant respects should get similar treatment. When this fails to happen—when racist or sexist policies are enacted, for instance—then something has gone wrong.
- *Self-interest isn't the only ethical consideration.* How well-off we are is important. But it isn't the only thing of moral importance. Morality sometimes calls on us to set aside our own interests for the sake of others.
- *Agony is bad.* Excruciating physical or emotional pain is bad. It may sometimes be appropriate to cause such extreme suffering, but doing so requires a very powerful justification.

- *Might doesn't make right.* People in power can get away with lots of things that the rest of us can't. That doesn't justify what they do. That a person can escape punishment is one thing—whether his actions are morally acceptable is another.
- *Free and informed requests prevent rights violations.* If, with eyes wide open and no one twisting your arm, you ask someone to do something for you, and she does it, then your rights have not been violated—even if you end up hurt as a result.

There are a number of points to make about these claims.

First, this short list isn't meant to be exhaustive. It could be made much longer.

Second, I am not claiming that the items on this list are beyond criticism. I am saying only that each one is very plausible. Hard thinking might weaken our confidence in some cases. The point, though, is that without such scrutiny, it is perfectly reasonable to begin our moral thinking with the items on this list.

Third, many of these claims require interpretation in order to apply them in a satisfying way. When we say, for instance, that equals ought to be treated equally, we leave all of the interesting questions open. (What makes people equals? Can we treat people equally without treating them in precisely the same way? Etc.)

Not only do we have a variety of plausible starting points for our ethical investigations; we also have a number of obviously poor beginnings for moral thinking. A morality that celebrates genocide, torture, treachery, sadism, hostility, and slavery is, depending on how you look at it, either no morality at all or a deeply failed one. Any morality worth the name will place *some* importance on justice, fairness, kindness, and reasonableness. Just how much importance, and just how to balance things in cases of conflict—that is where the real philosophy gets done.

## What Is Morality?

Before investing yourself too heavily in a subject matter, it would be nice to first have some idea of what you are getting yourself into. One way—sometimes the best—to gain such an understanding is by considering a definition. When you open your trigonometry text or chemistry handbook, you'll likely be given, very early on, a definition of the area you are about to intensively study. So, as a responsible author, I would seem to have a duty now to present you with a definition of *morality*.

I'd certainly like to. But I can't. There is no widely agreed-on definition of morality. We know that it is centrally concerned with protecting people's well-being, with fairness, justice, respect for others, virtue, responsibility, rights, liberties, social cooperation, praise, and blame. But the precise nature of such concern is highly disputed, as we'll soon see.

The absence of a definition does not leave us entirely in the dark. (After all, no one has yet been able to offer informative definitions of *literature*, or *life*, or *art*, and yet we know a great deal about those things.) We can get a sense of our subject matter by considering the questions that are distinctive of ethics—namely, those that structure this book, the ones focused on the good life, our duties to others, the nature of virtue, etc. We can supplement this by considering the starting points listed above, and many others that we could easily identify.

We can also better understand morality by contrasting its principles with those that govern the law, etiquette, self-interest, and tradition. Each of these represents a set of standards for how we ought to behave, ideals to aim for, rules that we should not break. But the fact that a law tells us to do something does not settle the question of whether morality gives its stamp of approval. Some immoral acts (like cheating on a spouse) are not illegal. And some illegal acts (like voicing criticism of a dictator) are not immoral. Certainly, many laws require what morality requires, and forbid what morality forbids. But the fit is hardly perfect, and that shows that morality is something different from the law. That a legislature passed a bill is not enough to show that the bill is morally acceptable.

We see the same imperfect fit when it comes to standards of etiquette. Forks are supposed to be set to the left of a plate, but it isn't immoral to set them on the right. Good manners are not the same thing as morally good conduct. Morality sometimes requires us *not* to be polite or gracious, as when someone threatens your children or happily tells you a racist joke. So the standards of etiquette can depart from those of morality.

The same is true when it comes to the standards of self-interest. I've just been watching the entire run of *The Shield*, a police drama set in a crime-ridden district of Los Angeles. Early in the series, the main character, Vic Mackey, murders a fellow police officer who was set to reveal Mackey's corruption. Mackey successfully frames a criminal for the murder: a classic case of protecting one's own interests by acting immorally. Though the relation between self-interest and morality is contested, it is a plausible starting point to assume that morality can sometimes require us to sacrifice our well-being, and that we can sometimes improve our lot

in life by acting unethically. So the standards of morality are not the very same as those of self-interest. (We will see a challenge to this view when considering *ethical egoism* in chapter 8.)

Finally, morality is also distinct from tradition. That a practice has been around a long time does not automatically make it moral. Morality sometimes requires a break with the past, as it did when people called for the abolition of slavery or for allowing women to vote.

True, people do sometimes speak of *conventional morality*, which is the set of traditional principles that are widely shared within a culture or society. These principles, like those of the law and etiquette, are the result of human decisions, agreements and practices. Conventional morality can differ from society to society. At least some of its principles can be traced to common misunderstandings, irrationality, bias, or superstition.

When I write about morality in this book, I am *not* referring to conventional morality. I am assuming that some social standards—even those that are long-standing and very popular—can be morally mistaken. (We'll examine this assumption at length in chapter 19.) So when I talk about morality from this point on, I will be referring to moral standards that are not rooted in widespread endorsement, but rather are independent of conventional morality and can be used to critically evaluate its merits.

It's possible, of course, that conventional morality is all there is. But this would be a very surprising discovery. Most of us assume, as I will do, that the popularity of a moral view is not a guarantee of its truth. We could be wrong on this point, but until we have a chance to consider the matter in detail, I think it best to assume that conventional morality can sometimes be mistaken. If so, then there may be some independent, "critical" morality that (i) does not have its origin in social agreements, (ii) is untainted by mistaken beliefs, irrationality, or popular prejudices, and (iii) can serve as the true standard for determining when conventional morality has got it right and when it has fallen into error. That is the morality whose nature we are going to explore in this book.

## Moral Reasoning

Moral reasoning, like all reasoning, involves at least two things: a set of reasons, and a conclusion that these reasons are meant to support. When you put these two things together, you have what philosophers call an **argument**. This isn't a matter of bickering or angrily exchanging words.

An argument is simply any chain of thought in which reasons (philosophers call these **premises**) are offered in support of a particular conclusion.

Not all arguments are equally good. This is as true in ethics as it is in science, mathematics, or politics. It is easy to mistake one's way when it comes to ethical thinking. We can land at the wrong conclusion (by endorsing child abuse, for instance). We can also arrive at the right conclusion by means of terrible reasoning. We must do our best to avoid both of these mistakes.

In other words, our moral thinking should have two complementary goals: getting it right, and being able to back up our views with flawless reasoning. We want the truth, both in the starting assumptions we bring to an issue and in the conclusions we eventually arrive at. But we also want to make sure that our views are supported by excellent reasons. And this provides two tests for good moral reasoning: first, we must avoid false beliefs, and second, the logic of our moral thinking must be rigorous and error-free.

The first test is pretty easy to understand. Consider the following quote, written in 1833, by the pro-slavery author Richard Colfax:

[T]he mind will be great in proportion to the size and figure of the brain: it is equally reasonable to suppose, that the acknowledged meanness of the negro's intellect, only coincides with the shape of his head; or in other words, that his want [i.e., lack] of capability to receive a complicated education renders it improper and impolitic, that he should be allowed the privileges of citizenship in an enlightened country.<sup>2</sup>

And here is William John Grayson, antebellum congressman and senator from South Carolina, on the same subject:

Slavery is the negro system of labor. He is lazy and improvident. . . . What more can be required of Slavery, in reference to the negro, than has been done? It has made him, from a savage, an orderly and efficient labourer. It supports him in comfort and peace. It restrains his vices. It improves his mind, morals and manners.<sup>3</sup>

---

2. Richard H. Colfax, *Evidence Against the Views of the Abolitionists, Consisting of Physical and Moral Proofs, of the Natural Inferiority of the Negroes* (New York: James T. M. Bleakley, 1833), p. 25.

3. William John Grayson, *The Hireling and the Slave* (Charleston, S.C.: John Russell, 1855), p. vii.

There are false beliefs galore in these (and other) defenses of American chattel slavery. Africans, and those of African descent, are not inherently lazy or unfit for a complicated education; they do not have heads with different shapes than whites; head shape is not correlated with intelligence; slavery is anything but comfortable. When one starts with false assumptions, the entire chain of reasoning becomes suspect. Good reasoning, in ethics as elsewhere, must avoid false beliefs if we are to have any confidence in its conclusions.

But it is possible to develop moral arguments that fail, even though every single one of their premises is true. The failure is of the second sort mentioned above: a failure of logic.

Consider this argument:

1. Heroin is a drug.
2. Selling heroin is illegal.
3. Therefore, heroin use is immoral.

This is a moral argument. It is a set of reasons designed to support a moral conclusion. Both of the premises are true. But they do not adequately support the conclusion, since one can accept them while consistently rejecting this conclusion. Perhaps the use of illegal drugs such as heroin really is immoral. But we need a further reason to think so—we would need, for instance, the additional claim that all drug use is immoral.

The argument in its present form is a poor one. But not because it relies on false claims. Rather, the argument's logical structure is to blame. The logic of an argument is a matter of how its premises are related to its conclusion. In the best arguments, the truth of the premises guarantees the truth of the conclusion. When an argument has this feature, it is **logically valid**.

The heroin argument is invalid. The truth of its premises does not guarantee the truth of its conclusion—indeed, the conclusion may be false.

Since the best arguments are logically valid, we will want to make sure that our own arguments meet this condition. But how can we do that? How can we tell a valid from an invalid argument, one that is logically perfect from one that is logically shaky?

There is a simple, three-part test:

1. Identify all of an argument's premises.
2. Imagine that all of them are true (even if you know that some are false).

3. Then ask yourself this question: Supposing that all of the premises were true, could the conclusion be false? *If yes*, the argument is invalid. The premises do not guarantee the conclusion. *If no*, the argument is valid. The premises offer perfect logical support for the conclusion.

Validity is a matter of how well an argument's premises support its conclusion. To test for this, we must assume that all of an argument's premises are true. We then ask whether the conclusion must therefore be true. If so, the argument is valid. If not, not.

Note that an argument's validity is a matter of the argument's structure. It has nothing to do with the *actual* truth or falsity of an argument's premises or conclusion. Indeed, *valid arguments may contain false premises and false conclusions*.

To help clarify the idea, consider the following argument. Suppose you are a bit shaky on your U.S. history, and I am trying to convince you that John Quincy Adams was the ninth president of the United States. I offer you the following line of reasoning:

1. John Quincy Adams was either the eighth or the ninth U.S. president.
2. John Quincy Adams was not the eighth U.S. president.
3. Therefore, John Quincy Adams was the ninth U.S. president.

In one way, this reasoning is impeccable. It is logically flawless. This is a valid argument. If all premises of this argument were true, then the conclusion would have to be true. It is impossible for 1 and 2 to be true and 3 to be false. It passes our test for logical validity with flying colors.

But the argument is still a bad one—not because of any logical error, but because it has a false premise (number 1; Quincy Adams was the sixth U.S. president) and a false conclusion. The truth of an argument's premises is one thing; its logical status is another.

The lesson here is that truth isn't everything; neither is logic. We need them both. What we want in philosophy, as in all other areas of inquiry, are arguments that have two features: (1) they are logically watertight (valid), and (2) all of their premises are true. These arguments are known as **sound** arguments.

Sound arguments are the gold standard of good reasoning. And it's easy to see why. They are logically valid. So if all of their premises are true, their conclusion must be true as well. And by definition, sound arguments contain only true premises. So their conclusions are true. If you can tell

that an argument is valid, and also know that each premise is correct, then you can also know that the conclusion is true. That is what we are after.

I started this section by claiming that not all moral arguments are equally good. We're now in a position to see why. Some arguments rely on false premises. Others rely on invalid reasoning. Still others—the worst of the lot—commit both kinds of error.

To reinforce these points, consider one more moral argument. Some people say that killing animals and eating meat is morally okay, because animals kill other animals, and there is nothing immoral about that. Is this a plausible line of reasoning?

Not as it stands. To see this, let's reconstruct the argument by stating it in premise-conclusion form. This is something that I'm going to do for dozens of arguments over the coming pages. For those of you who want to improve your philosophy skills, there's no better way to do so than to take a line of reasoning in ordinary English and try to set it out step by step. That makes it easier to tell just what is being claimed, and so easier to determine the logical structure of the argument and the truth of the premises.

Here is my take on this popular *Argument for Meat Eating*:

1. It is morally acceptable for non-human animals to kill and eat other animals.
2. Therefore, it is morally acceptable for human beings to kill and eat non-human animals.

As stated, there is only one premise to this argument. And it is true. So if the argument is problematic, it has to be because of its logic.

And that is indeed its flaw. The argument is invalid; the premise does not adequately support the conclusion. We can assume that the premise is true (indeed, we *should* accept it), but the conclusion might still be false. The truth of the premise is not enough to guarantee the truth of the conclusion, since what is morally acceptable for animals may not be morally acceptable for us. We would need a further premise, to the effect that we are allowed to do anything that animals do, in order to make this argument valid.

So as it stands, the Argument for Meat Eating is invalid. Therefore it is unsound. Does that mean that its conclusion is false?

No. And here is another important lesson about reasoning: bad arguments may contain true conclusions. After all, even true claims can be supported by poor reasoning. The fact that the Argument for Meat Eating

is invalid does not show that its conclusion is false. It shows only that this particular way of defending that conclusion is no good. For all we know, there might be other, better arguments that can do the trick.

The Argument for Meat Eating, like many other invalid arguments, can be modified so that it takes on a logically perfect form. Indeed, a charitable reading of the argument would show that there is an underlying assumption that, if brought out into the open, would allow us to transform the argument into a valid one. With a little tweaking, for instance, we get the following:

1. If it is morally acceptable for non-human animals to kill and eat one another, then it is morally acceptable for humans to kill and eat non-human animals. (This is the underlying assumption.)
2. It is morally acceptable for non-human animals to kill and eat one another.
3. Therefore, it is morally acceptable for humans to kill and eat non-human animals.

And this argument is logically perfect. If premises 1 and 2 are true, then the conclusion, 3, has to be true.

But even this version is unsound. Not because it is invalid, but because it now contains a false premise. Premise 2 is true. But premise 1 is not. Four reasons explain this.

First, animals that eat other animals have no choice in the matter. We do.

Second, a carnivore's survival depends on its eating other animals. Ours does not. With rare exceptions, human beings can survive perfectly well without eating animal flesh. There are hundreds of millions of vegetarians leading healthy lives.

Third, none of the animals we routinely eat (chickens, cows, pigs, sheep, ducks, rabbits) are carnivores. They *don't* eat other animals. So if their behavior is supposed to guide our own, then we should eat only plants.

Fourth, it is implausible to look to animals for moral guidance. Animals are not **moral agents**—they can't control their behavior through moral reasoning. That explains why they have no moral duties, and why they are immune from moral criticism. But we, obviously, are moral agents, and we can guide our behavior by the moral decisions we make.

Again, this analysis does not prove that the argument's conclusion is false. It just shows that this version of the argument, like the original, is unsound. Meat eating may be perfectly morally acceptable. But this argument fails to show that.

I have spent some time on this argument, not because I want to defend a view about whether vegetarianism is morally required, but because I want to illustrate the possibility of real moral argumentation. We started with a version of the argument that has convinced a lot of people. But when we laid it out clearly, we could see that it was invalid. So we modified it, making an underlying assumption explicit, and doing so in a way that gave us a logically perfect argument. But even this improved version is unsound, because its first premise is false.

Can I be absolutely sure that the premise is false? No. I will be the first to admit that further argument might reveal the error of my thinking. What's more, there is no foolproof method that can perfectly sort true claims from false ones. We may offer excellent reasons and arguments on behalf of our moral views, but at the end of the day, it's possible that not everyone will be convinced.

But this is no different from any other area of inquiry. There is no litmus test that can distinguish all true biological claims from false ones, accurate economic forecasts from the inaccurate, correct chemistry hypotheses from incorrect ones. There is potential for disagreement in all areas of thinking.

The absence of a perfectly reliable test for truth does not mean that all claims are equally true, or that truth is in the eye of the beholder. The earth is not a cube. Six is less than ten. Queen Victoria is dead. Cats are animals. These claims are each true. Their opposites are false. And our say-so has nothing to do with it. These claims would be true even if we were not around to make them. They aren't true because we think they are; we think they are true because they are.

Perhaps things are this way in ethics, too. We will spend a lot of time considering whether that is so when we discuss metaethics in the last part of the book. For right now, the important thing to note is that we must rely on our good sense and good judgment in all areas of investigation, not just in ethics. The lack of a precision test for truth does not spell the defeat of moral inquiry, since other areas of investigation get along just fine without one.

Moral reasoning is just what its name implies—offering and evaluating reasons that are meant to support moral conclusions. It is not merely a matter of doing a gut check and venting one's feelings. Not every reason is a good one. Some reasons fail to support their conclusions. Others represent false beliefs. And while it is sometimes hard to separate fact from fiction, this needn't hobble us. Many claims are clearly true, many clearly

false. For the others, there is evidence and argument that we can bring in to try to settle the matter. This won't always yield decisive results. But that's the nature of our situation. We can't always be sure of things, in ethics or elsewhere. That shouldn't prevent us from trying to get it right, and backing up our moral views with the best possible reasons.

## The Role of Moral Theory

A great deal of philosophy is done at a pretty high level of abstraction. That's not necessarily a bad thing, even though reading and thinking at that level is typically more challenging and less fun than getting engrossed in the details of a well-written novel or historical narrative. Of course we'll need to get back down to earth and familiarize ourselves with the specific facts of a case before knowing what to do in a given situation. But according to most philosophers, knowing what to do here and now also requires that we have a sure grasp of very general moral principles. Knowing which principles are plausible, and how they relate to one another, is a large part of what moral philosophy is all about.

Moral philosophy is primarily a matter of thinking about the attractions of various ethical *theories*. When we develop and test these theories, we are bound to look beyond the details of specific cases. We are trying to find the deepest truths about our subject matter—how to live. Such truths are wide-ranging and apply to countless cases. That's why moral philosophers so often look beyond the details of specific cases and focus instead on very general principles.

Moral theorizing is the result of a perfectly natural process of thinking. We are questioning beings, interested in seeking out ever-deeper explanations of things. And we are uneasy if there is no chance of a unifying explanation, an account that can coherently organize the various aspects of our thinking and experience. This is clear in psychology, for instance, where researchers have always been drawn to unifying views of human motivation. For many psychologists, it all comes down to self-interest (egoists), or to how we have been conditioned (behaviorists), or our sexual impulses (Freudians), etc. This process is evident in physics, too, where the dream is one day to discover the unified theory—a single master principle that will explain *all* of the workings of the physical world, from the movements of subatomic particles to the behavior of the largest stars and galaxies.

The same desire for unification is also present in ethics. We might begin a conversation by insisting on the immorality of some specific

action—say, revealing a patient's confidential information. But someone might challenge our view, and in reply, we would cite a moral rule to back it up: revealing such information is wrong because it betrays a trust. But why is it wrong to betray a trust? Because (we might say) such actions fail to show respect for the person who has been betrayed. But why is it wrong to fail to show respect? And is it always wrong to do such a thing, or are there exceptions? If there are exceptions, what explains *them*? This is a perfectly natural way of going on. We are searching for increasingly general moral principles with the power to explain more and more cases. The hope is eventually to land on just a single principle, one that will do all of the explaining we need in the moral realm.

Suppose that we think really carefully about our moral beliefs, and find that we ultimately justify them by means of four principles:

- Don't impose unnecessary harm.
- Be nice to others.
- Act justly.
- Tell the truth.

Is there a next step? Of course! Aren't you curious to know whether there is a yet more general rule, one that can explain why these four principles are justified? Like researchers in most areas, moral philosophers remain dissatisfied unless they can offer a truly comprehensive theory that will unify and impose order on our thoughts. Physicists want this. Psychologists want this. So do philosophers.

That's why our focus will mostly be on these very general ethical theories. They represent the natural outgrowth of some extremely compelling ethical ideas—ones that you surely have relied on when trying to justify your own moral views. There is something important in taking our core ethical beliefs and seeing where they lead. They lead to ethical theories; doing moral philosophy is the process of tracing the lines that connect our basic moral views to these more developed theories, and then testing them to see how well they can hold up against our curiosity and critical intelligence.

## Looking Ahead

In the pages to come, I present and evaluate a lot of arguments. These are the ones at the very heart of morality, the ones that try to offer answers to the deepest questions of ethics. As we will see, no fundamental

theory—about the good life, our moral duties, or the status of morality—has earned anything like unanimous support among philosophers.

I say this not to dash your hopes, but to give you a realistic take on what to expect. There is a very broad consensus on a number of points in ethics. Consider, for instance, the twelve starting points mentioned earlier in this introduction, a sampling that could easily have been expanded. The moral issues that tend to capture our attention are those that are hotly disputed. What often goes unnoticed is the substantial amount of moral agreement, even across time and place.

Still, when it comes to devising a theory that can offer a comprehensive account of morality, things become much trickier. And then a natural, despairing thought: greater minds than ours have spent lifetimes trying to solve the core questions of ethics, and none of their theories has gained universal support. So what's the use?

It's a fair question. But there is a good answer. We are thinking about how to live; what could be more important than that? We can make a lot of progress in our own thinking by studying the thoughts and arguments of those who have devoted so much effort to this vital task. We may realize that our own “philosophy of life” is marred in ways that we hadn't foreseen. Or we might come to appreciate certain benefits of our views that had escaped our notice. Those of you who work your way through this book will certainly be in a much better position to critically assess your own moral views, and to improve your thinking about how to live your life.

What ethicists across the ages have done is to take a common insight—say, that happiness is the key to a good life, that we must treat everyone fairly, that we must prevent harm—and see how far we can get by consistently applying this insight. Consistency is not to be sneezed at. It's not the hobgoblin of little minds, but a minimum test of a theory's plausibility. Inconsistent, contradictory views cannot be true, which is why philosophers try so hard to avoid them.

Suppose that you are involved in a moral debate, or are thinking about how to improve your own life. If you go deep enough, you'll probably land on a view that you can no longer defend. Perhaps it's one of the twelve mentioned earlier. Perhaps it's something else. Whatever it is, the truth of that view is important. And unsurprisingly, philosophers across the ages will have examined that view very carefully. We can learn from their work. We can find out what is attractive about these starting points. And we can also discover how they might be vulnerable.

That's not everything. Agreed. You won't find, by the end of the book, a recipe for the best life, or a simple step-by-step guide for doing your duty. This book does not belong on the self-help shelves. You probably already figured that out, since such manuals are a lot chattier and far easier to read than this one. But those books never get to the deepest issues—most of them assume, for instance, that happiness is what we should be trying for, or that getting what you want is what life is all about. Philosophers subject such thoughts to intense scrutiny. And it isn't clear whether they survive.

Let's start our work together by having a look at these views, ones that focus on the good life for human beings. There are a lot of surprises in store.

## Discussion Questions

1. What are the three main areas of moral philosophy? How might answering questions in one area inform the way we think about the other areas?
2. What is moral skepticism, and what are some reasons that people give for being moral skeptics? Do you think moral skepticism is correct? Why or why not?
3. Do you find the “ethical starting points” presented here to be plausible? Do you think we can prove that such starting points are correct? Do we need to?
4. How do moral standards differ from those of the law, etiquette, or tradition?
5. What are the two ways that a moral argument can fail? Can an argument with either of these failings still have a true conclusion? Why or why not?
6. Is it possible for a valid argument to be unsound? What about for a sound argument to be invalid? For each, either give an example of such an argument or explain why it is impossible.
7. What is the point of trying to develop a moral theory? Do you think that there might be a single, ultimate moral principle that can explain the truth of all other moral principles?